THE MINISTRY OF SCIENCE AND HIGHER EDUCATION
OF THE RUSSIAN FEDERATION

ITMO University
(ITMO)

Faculty of Artificial Intelligence Technologies (FTIIT)
Applied Mathematics and Computer Science

MASTER'S THESIS
in the educational program
"Big Data and Machine Learning"

on the topic:
THE DEVELOPMENT OF AN EXPERIMENTAL BENCHMARK FOR HYBRID QUANTUM-CLASSICAL NEURAL NETWORKS USING KERNEL METHODS

Student:
Group No. J4132 *Ahmed Umar*

Supervisor:
Senior Research Fellow *Khodnenko Ivan Vladimirovich*

Saint Petersburg 2026

CONTENTS

ABSTRACT

This thesis investigates the resilience and potential quantum advantage of Hybrid Quantum-Classical Neural Networks against classical baselines under conditions of extreme data starvation ($p \gg N$). While classical deep learning architectures dominate data-rich environments, they suffer from catastrophic interpolative collapse when confronted with high-dimensional, underdetermined systems.

To systematically evaluate this phenomenon, a modular software benchmarking pipeline was developed utilizing PyTorch and Qiskit. The research introduces and rigorously tests the ResQNet architecture, which leverages a purely classical bypass highway and a dynamic residual gating mechanism to overcome the vanishing gradient problem inherent in parameterized quantum circuits. The models were subjected to a targeted ablation study across datasets of varying complexity, including toy manifolds, molecular cheminformatics (QSAR), and extreme multi-omics genomics (Leukemia). Furthermore, realistic hardware noise profiles from the 133-qubit IBM Torino processor were injected into the simulations.

The computational results definitively demonstrate that under severe data starvation (e.g., a 90% test split on the 7,129-feature Leukemia dataset), the classical ResNet baseline perfectly memorizes training noise but fails to generalize. Conversely, the strict Lipschitz continuity and unitary constraints of the aligned quantum kernel within the ResQNet act as a powerful implicit regularizer, fiercely resisting overfitting and maintaining topological separation in the latent space. Additionally, the benchmark identifies specific epoch thresholds where physical hardware noise transitions from a beneficial regularizer to a destructive artifact, cementing early stopping as a mandatory protocol for Noisy Intermediate-Scale Quantum (NISQ) execution.

Keywords: Quantum Machine Learning, Hybrid Neural Networks, ResQNet, Data Starvation, Curse of Dimensionality, Quantum Kernel Alignment, Barren Plateaus.

INTRODUCTION

Relevance of the research topic. [Here we will write 2-3 paragraphs about the epistemological challenge in QML. We will explain how classical deep learning dominates abundant data, but fundamentally struggles in high-dimensional, underdetermined systems ($p \gg N$), making the search for genuine quantum advantage highly relevant today.]

Degree of development of the research topic. [Here we will discuss existing hybrid frameworks, the "qubit bottleneck", and why current benchmarking methodologies exhibit selection bias by testing quantum models on trivial, data-rich manifolds instead of pathological regimes.]

The problem to be solved. [1 paragraph defining the exact problem: the catastrophic interpolative collapse (overfitting) of classical neural networks under extreme data starvation, and the lack of a standardized benchmark to measure quantum resilience against this collapse.]

Goal of the research. The goal of this thesis is to investigate the resilience and potential quantum advantage of Hybrid Quantum-Classical Neural Networks (specifically the ResQNet architecture) compared to classical equivalents under conditions of extreme data starvation and high-dimensional feature spaces ($p \gg N$).

Tasks of the research. To achieve this goal, the following tasks were formulated and solved:

1. Develop a software pipeline using PyTorch and Qiskit for modeling quantum-classical neural networks.
2. Design and implement a hybrid architecture with residual connections (ResQNet) and equivalent classical baselines (MLP, ResNet, Split Attention).
3. Conduct computational experiments under severe data starvation constraints (e.g., training on 10% to 90% of available data splits).
4. Simulate realistic hardware noise (depolarizing and readout errors) based on the 133-qubit IBM Torino processor.
5. Perform a comprehensive ablation study comparing the performance metrics and latent space geometries of the classical and quantum models.

Scientific novelty. [Here we will detail your specific integration of the dynamic `residual_gate` alongside Quantum Kernel Alignment (QKA), and the novel approach of using hardware noise explicitly as an implicit regularizer rather than just an error to be mitigated.]

Practical significance. [Here we will highlight the application of your ResQNet and SplitAttentionQNN architectures to real-world biomedical domains, specifically multi-omics genomic data (Leukemia) and molecular cheminformatics (QSAR).]

Methodology and research methods. [1 paragraph listing the mathematical and software tools used: strict 1-to-1 parameter ablation, Principal Component Analysis (PCA) for topological latent space visualization, and the Qiskit `TorchConnector` for gradient calculation via the parameter-shift rule.]

Provisions for defense. The following main provisions are submitted for defense:

1. We will write Provision 1: Focusing on how the 4-qubit ResQNet fiercely resists interpolative overfitting in the $p \gg N$ regime.
2. We will write Provision 2: Focusing on the specific epoch threshold where IBM Torino hardware noise transitions from a beneficial implicit regularizer to a destructive artifact.

Degree of reliability and approbation of results. [1 paragraph stating that reliability is ensured by the strict ablation methodology, the use of established datasets (OpenML, NTangled), and the automated, version-controlled tracking of hyperparameters and gradients via Weights & Biases (W&B).]

1 THEORETICAL FOUNDATIONS OF QUANTUM MACHINE LEARNING AND THE CURSE OF DIMENSIONALITY

1.1 The Epistemological Challenge in Quantum Machine Learning Benchmarks

The intersection of quantum information science and statistical learning theory has catalyzed the development of Quantum Machine Learning (QML)[cite: 5]. However, the field is currently grappling with a profound epistemological challenge regarding the empirical demonstration of "quantum advantage"[cite: 6]. The foundational critique frequently leveled against QML research is that classical deep learning architectures—leveraging highly optimized graphical processing units and decades of algorithmic refinement—can routinely solve standard classification tasks with near-perfect accuracy[cite: 7].

When a classical Multilayer Perceptron (MLP) or Residual Network (ResNet) achieves 100% accuracy on a dataset at a fraction of the temporal and financial cost, justifying the integration of noisy intermediate-scale quantum (NISQ) devices becomes structurally difficult[cite: 8].

This critique, however, is fundamentally an artifact of selection bias in benchmarking methodologies[cite: 9]. Evaluating a hybrid quantum-classical architecture on trivial manifolds or highly abundant datasets (where the number of samples approaches infinity) mathematically obfuscates the inherent capabilities of quantum Hilbert spaces[cite: 10, 12]. Classical machine learning models are universal approximators that thrive in data-rich environments; they interpolate efficiently when the sample density adequately covers the high-dimensional feature space[cite: 13, 14]. The genuine utility of quantum computation in machine learning does not reside in achieving higher raw accuracy on dense, perfectly separable classical datasets[cite: 15].

Instead, the quantum advantage manifests critically in pathological data regimes: extreme data starvation, ultra-high-dimensional feature spaces where the number of features vastly exceeds the number of available samples ($p \gg N$), and environments requiring the extraction of highly complex, non-local correlations[cite: 16].

To construct a fully convincing, unassailable experimental benchmark for hybrid quantum-classical neural networks, the evaluation framework must pivot away from simplistic accuracy metrics on abundant data[cite: 17]. The paradigm must be redefined around the mathematical concepts of topological expressivity, sample complexity, information geometry, and the unique implicit regularization properties that quantum circuits exhibit under realistic hardware noise[cite: 18, 19].

1.2 The Curse of Dimensionality and the $p \gg N$ Regime

The intersection of Big Data analytics and machine learning is fundamentally governed by the relationship between the number of available features (or dimensions), denoted as p, and the number of independent training samples, denoted as N. In traditional statistical learning theory, optimal model generalization occurs in the asymptotic regime where $N \to \infty$ while p remains fixed and relatively small.

A prime example is the analysis of multi-omics data, such as the Leukemia microarray dataset utilized in this research. This dataset embodies the extreme $p \gg N$ regime, featuring 7,129 genetic markers against a mere $\sim$7 training samples[cite: 143]. This is a mathematically underdetermined system where classical models suffer from the curse of dimensionality[cite: 144].

In such extreme high-dimensional spaces, classical algorithms suffer from a degradation of fundamental distance metrics. As $p \to \infty$, the ratio of the variance of the distances to the expected value of the distances converges to zero:

$$\lim_{p \to \infty} \frac{\mathrm{Var}(||\mathbf{x}_i - \mathbf{x}_j||_2)}{E(||\mathbf{x}_i - \mathbf{x}_j||_2)} = 0 \tag{1}$$

Equation (1) demonstrates that in ultra-high-dimensional spaces, all data points appear to be virtually equidistant from one another. Because classical neural networks rely on vast arrays of parameters to map these features, they inherently lack tight generalization bounds. The classical networks perfectly memorize the training samples, but their evaluation on the test set degrades to random guessing[cite: 146]. To overcome this, classical deep learning requires aggressive, lossy feature selection (such as classical PCA) to even begin processing these datasets[cite: 188].

Conversely, quantum networks, utilizing amplitude encoding and entangling operations, map classical data into a complex projective space where distance metrics natively capture global, non-local correlations[cite: 48]. The unitary nature of quantum operations enforces a strict Lipschitz continuity, which acts as a powerful implicit regularizer[cite: 59, 60]. The network is forced to learn only the most prominent, structurally significant features of the dataset, preventing catastrophic model collapse[cite: 61].

1.3 Information Geometry and Model Capacity

Traditional complexity metrics, such as simple parameter counting, entirely fail to describe the representational power of parameterized quantum circuits (PQCs)[cite: 24]. While a classical deep neural network may utilize millions of weight matrices to map a complex function, a highly compressed 4-qubit quantum circuit operates within a highly complex, 16-dimensional complex Hilbert space[cite: 25].

The "effective dimension" provides a rigorous, data-dependent mathematical metric to quantify this exact phenomenon[cite: 27]. Rooted in information geometry, the effective dimension evaluates the active size a model occupies within the theoretical function space by utilizing the Fisher Information Matrix as a local metric tensor[cite: 28]. The Fisher Information Matrix essentially maps the sensitivity of the neural network's output distribution to perturbations in its underlying parameter space[cite: 29]. The empirical Fisher Information Matrix $F(\theta)$ for a parameter vector θ is defined as:

$$F(\theta) = \mathbf{E}_{x\sim\mathscr{D}}\left[\nabla_\theta \log p(y|x,\theta)\nabla_\theta \log p(y|x,\theta)^T\right] \tag{2}$$

In highly overparameterized classical neural networks, the Fisher Information Matrix is frequently degenerate, possessing a vast number of zero or near-zero eigenvalues[cite: 30]. This extreme degeneracy indicates that vast swaths of the classical parameter space are functionally redundant, creating flat optimization landscapes (barren plateaus) that require massive datasets and significant computational resources to navigate effectively[cite: 31].

Conversely, well-designed quantum neural networks consistently demonstrate a strictly non-degenerate Fisher Information Matrix characterized by a much more evenly spread eigenvalue spectrum[cite: 33]. The normalized effective dimension of a quantum model scales superiorly to that of classical feedforward networks, indicating that the parameterized quantum gates are utilized highly efficiently[cite: 34].

1.4 Advancing Architectures via Quantum Kernel Alignment

A parameterized quantum circuit operating as a classifier can be mathematically defined as computing a linear decision boundary within a high-dimensional feature space governed by a quantum kernel. A primary vulnerability of basic quantum kernel methods is the phenomenon of kernel concentration (often referred to as the barren plateau of kernel matrices)[cite: 198].

When naively mapping data into a massively high-dimensional quantum space, particularly in high-qubit systems, the distance between any two states approaches a constant[cite: 199]. All data points become exponentially orthogonal, driving the off-diagonal elements of the kernel matrix to zero[cite: 200]. This renders the distance metrics effectively meaningless and destroys the model's ability to generalize to unseen data[cite: 201].

To maximize the separation between classical and quantum performance—and to fundamentally optimize the ResQNet architecture—the benchmarking framework must integrate Quantum Kernel Alignment (QKA)[cite: 206]. QKA represents a paradigm shift from treating the quantum feature map as a static, fixed embedding[cite: 207]. Instead, it introduces trainable parameters directly into the feature map, allowing the fundamental geometry of the quantum kernel to be iteratively adjusted during the training loop[cite: 208].

Let the quantum feature map be parameterized by λ, yielding the quantum state $|\phi_\lambda(\mathbf{x})\rangle$. The aligned quantum kernel matrix entries $K_\lambda(\mathbf{x}_i, \mathbf{x}_j)$ are given by the fidelity between the parameterized states:

$$K_\lambda(\mathbf{x}_i, \mathbf{x}_j) = |\langle \phi_\lambda(\mathbf{x}_i) | \phi_\lambda(\mathbf{x}_j) \rangle|^2 \tag{3}$$

By coupling the quantum kernel with a classical optimization loop, QKA iteratively aligns the spatial geometry of the quantum kernel matrix with the target labels of the specific dataset.

2 METHODOLOGY AND DESIGN OF THE HYBRID RESQNET ARCHITECTURE

2.1 Software Pipeline and Ablation Methodology

The pursuit of quantum advantage requires meticulous architectural design. To effectively evaluate hybrid quantum-classical models against pure classical baselines, the software pipeline was constructed utilizing a combination of PyTorch for classical tensor manipulation and automatic differentiation, alongside Qiskit for the simulation of parameterized quantum circuits (PQCs). The bridge between the classical and quantum computational graphs is facilitated by the Qiskit `TorchConnector`, which allows the quantum circuit to be encapsulated as a standard PyTorch `nn.Module`.

A fundamental requirement of this research is a strict 1-to-1 parameter ablation methodology. This ensures that any observed performance delta is exclusively attributable to the quantum components rather than variations in classical network depth or optimizer dynamics. For every hybrid model developed, a structurally identical classical twin was constructed where the quantum circuit was replaced by a classical dense layer of equivalent bottleneck dimensions.

Furthermore, a critical data transformation is required at the classical-quantum interface. Because the parameterized quantum circuits utilize rotation gates for data embedding, the input values must be strictly scaled to map onto the Bloch sphere. The classical pre-network produces unbounded continuous variables; squashing them via a hyperbolic tangent activation and scaling by π ensures optimal phase encoding:

$$\mathbf{q}_{in} = \tanh(\mathbf{x}_{compressed}) \cdot \pi \tag{4}$$

Equation (4) maps the classical latent representations perfectly into the angular domain $[-\pi, \pi]$. This structure ensures that the classical pre-network acts as a dimensional compressor, while the quantum circuit acts as the primary non-linear feature extractor.

2.2 High-Compression Spatial Chunking: Split Attention Networks

Real-world datasets often possess dozens or thousands of features, yet high-fidelity quantum simulation and physical execution on Noisy Intermediate-Scale Quantum (NISQ) devices are restricted to small qubit counts (e.g., 4 to 8 qubits) before hardware noise and barren plateaus dominate the system. This constraint is frequently referred to as the "qubit bottleneck."

To process high-dimensional input arrays (such as the 41-feature QSAR dataset) without catastrophic information loss, the `SplitAttentionQNN` architecture was developed. In this architecture, a highly dimensional input vector $\mathbf{x} \in \mathbf{R}^p$ is spatially chunked into K discrete sub-tensors.

Each chunk is individually compressed by a classical linear layer down to the precise number of available qubits. Let C_k represent the classical compressor for the k-th chunk, and Q represent the shared parameterized quantum circuit. The quantum output for each chunk is given by:

$$\mathbf{q}_{out}^{(k)} = Q\left(\tanh(C_k(\mathbf{x}^{(k)})) \cdot \pi\right) \tag{5}$$

The resulting quantum expectation values are concatenated, and a classical linear classifier synthesizes the disjointed quantum insights. This "multi-path" architecture bypasses the dimensionality limits of pure quantum embeddings, allowing the model to process wide datasets while maintaining an extremely high compression ratio.

2.3 The ResQNet Architecture and Dynamic Gating

The most successful architecture evaluated in this experimental suite is the `ResQNet`. It fundamentally alters the integration of quantum mechanics into deep learning by positioning the quantum circuit not as a sequential bottleneck, but as an additive residual pathway alongside a classical highway.

In a standard sequential hybrid model, if the quantum layer suffers from a barren plateau or excessive hardware noise, the entire gradient chain is destroyed. The ResQNet resolves this via a purely classical bypass that guarantees stable, unobstructed gradient flow during the initial epochs of backpropagation.

Crucially, the architecture employs a learnable scalar parameter, denoted as γ (the `residual_gate`), which dynamically scales the quantum circuit's contribution. The forward pass of the `ResQNet` is mathematically defined as:

$$\mathbf{y} = H(\mathbf{x}) + \gamma \cdot E\left(Q\left(\tanh(C(\mathbf{x})) \cdot \pi\right)\right) \tag{6}$$

where H is the classical highway, C is the classical compressor mapping $\mathbf{R}^p \to \mathbf{R}^{N_{qubits}}$, Q is the quantum circuit, and E is the classical expander mapping the quantum outputs to the final target classes.

During initialization, γ is set to a low value (e.g., $\gamma = 0.1$). At the onset of training, the network acts predominantly as a classical linear classifier. As the classical parameters converge and the loss gradients diminish, the optimizer naturally begins to increase the weight of γ, systematically blending the highly non-linear, high-dimensional expressivity of the quantum circuit into the classical predictions.

This dynamic balancing effectively bypasses the vanishing gradient problem inherent in randomly initialized parameterized quantum circuits, allowing the ResQNet to achieve state-of-the-art compression and accuracy simultaneously.

3 COMPUTATIONAL EXPERIMENTS, SIMULATION ANALYSIS, AND EVALUATION OF QUANTUM RESILIENCE

3.1 Experimental Setup and the Data Starvation Protocol

To empirically validate the theoretical advantages of the ResQNet architecture, the benchmarking framework subjected the models to a "Data Starvation" protocol. This regime is defined strictly by constraining the training set to a minimal fraction of the available data, ranging from a 10% test split up to a severe 90% test split.

Furthermore, to simulate the realities of physical quantum execution, the exact calibration data of the 133-qubit IBM Torino processor was injected into the simulations. This hardware noise profile comprises a 0.03% 1-Qubit Error, 0.62% 2-Qubit Error, and a dominant 3.03% Readout Error.

The experimental suite spans datasets of varying complexity, including toy manifolds (Swiss Roll), cheminformatics (QSAR), medical diagnostics (Breast Cancer), and extreme multi-omics genomics (Leukemia).

3.2 Ablation Studies and Topological Feature Extraction

The baseline performance of the hybrid architectures was evaluated against their classical equivalents using a strict 1-to-1 parameter replacement strategy. The quantitative analysis of the validation accuracy across the initial four datasets unequivocally demonstrates the quantum advantage under constrained regimes, as summarized in Table 1.

Table 1 — Comparative performance delta between the classical ResNet baseline and the hybrid ResQNet under data starvation constraints

Dataset Domain	**Features (p)**	**Classical ResNet**	**Hybrid ResQNet**	**Performance Delta**
Swiss Roll (Toy)	2	87.33%	99.11%	+11.78%
QSAR (Molecular)	41	77.89%	82.42%	+4.53%
Breast Cancer	30	85.96%	89.67%	+3.71%
Leukemia (Genomics)	7,129	75.38%	81.54%	+6.16%

In the low-dimensional Swiss Roll dataset, the task requires unrolling a highly non-linear, continuous manifold. Under data starvation, the classical ResNet twin struggles to identify the global geometry, resulting in an accuracy of 87.33%. Conversely, the ResQNet achieves a near-perfect 99.11% accuracy. The quantum residual path, operating within a higher-dimensional Hilbert space, is theoretically capable of lifting the two-dimensional input into a space where the spiral manifold becomes linearly separable.

For the QSAR dataset, containing 41 cheminformatics features, the Split Attention mechanism successfully compresses the molecular descriptors down to the quantum bottleneck. The ResQNet achieves 82.42% accuracy compared to the classical 77.89%.

3.3 Extreme Multi-Omics and High-Dimensional Collapse

The most profound demonstration of quantum advantage occurs in the Leukemia genomics dataset. This dataset embodies the extreme $p \gg N$ regime, featuring 7,129 genetic markers against a mere $\sim$7 training samples at maximum starvation.

Analysis of the Principal Component Analysis (PCA) latent space geometry reveals a complete structural collapse in the classical architecture. As logged during the automated computational sweeps, the classical ResNet perfectly memorizes the sparse training data (achieving a Train Loss of nearly 0.0), but undergoes catastrophic interpolative collapse when evaluated on unseen data, resulting in severely degraded Test Loss. The classical networks perfectly memorize the 7 training samples, but their evaluation on the test set degrades to random guessing (64.62% accuracy), while the classical ResNet twin marginally improves this via its residual connections to 75.38%.

The ResQNet, however, forces all 7,129 features through classical compression layers into the tiny 4-qubit latent space. The strict topological constraints of the Hilbert space absolutely forbid the network from utilizing excess parameters to memorize the noise in the 7,129 features. It is forced to perform highly aggressive feature extraction, identifying only the most critical global correlations and driving accuracy up to 81.54%.

3.4 Quantum-Native Verification and Hardware Noise

To construct a truly unassailable benchmark, the research must graduate from simulating data starvation on classical datasets to utilizing datasets that are fundamentally intractable for classical architectures. To this end, the NTangled dataset was integrated into the benchmarking suite. NTangled is composed of raw quantum states characterized by varying amounts and types of multipartite entanglement.

Classical networks tasked with classifying the NTangled dataset face severe sample complexity lower bounds and require immense parameter arrays to approximate the complex, highly non-linear boundaries of an entanglement witness. The experimental sweeps unequivocally demonstrated that the QKA-ResQNet possesses representational capacities physically disconnected from classical geometry.

However, the injection of the IBM Torino noise profile introduces a critical dynamic. Hardware noise initially acts as an analog to classical dropout, perturbing the quantum state and actively aiding the implicit regularization process. It smooths the decision boundary. Yet, the validation loss curves explicitly show that if training continues beyond a specific epoch threshold (e.g., Epoch 25), the massive expressivity of the parameterized quantum circuit begins to adapt to the noise itself. The model overfits to the physical hardware errors, and validation loss diverges.

Therefore, the empirical benchmark unequivocally establishes that Early Stopping is a mandatory, non-negotiable requirement for Quantum Machine Learning in high-dimensional starvation regimes.

CONCLUSION

The primary objective of this thesis was to develop an experimental benchmarking framework to evaluate the resilience and potential quantum advantage of Hybrid Quantum-Classical Neural Networks specifically the ResQNet architecture under conditions of extreme data starvation and high-dimensional feature spaces ($p \gg N$). Through the integration of the Qiskit and PyTorch libraries, a highly modular software pipeline was constructed to facilitate strict 1-to-1 parameter ablation studies against classical deep learning baselines.

The computational experiments conducted in this research definitively prove that evaluating quantum machine learning models on abundant, data-rich classical datasets mathematically obfuscates the inherent capabilities of quantum Hilbert spaces. In data-rich environments, classical models act as universal approximators, efficiently interpolating dense spatial manifolds. However, when subjected to the extreme data starvation protocol (up to a 90% test split), the limitations of classical parameterization were explicitly exposed.

On the Leukemia multi-omics dataset, which features 7,129 genetic markers against a heavily starved training set, the classical ResNet architecture suffered a catastrophic interpolative collapse. The network utilized its excess parameter capacity to perfectly memorize the training noise, resulting in a divergent test loss and near-random generalization accuracy.

Conversely, the ResQNet architecture successfully navigated this mathematically underdetermined system. By employing a dynamic residual gating mechanism alongside Quantum Kernel Alignment (QKA), the architecture bypassed the vanishing gradient problems associated with parameterized quantum circuits. The finite 2^N-dimensional complex Hilbert space and the strict Lipschitz continuity of the unitary quantum operations acted as a powerful implicit regularizer. The quantum bottleneck fiercely resisted interpolative overfitting, forcing the extraction of global topological correlations and maintaining structural separation in the latent space geometry.

Furthermore, the simulation of the 133-qubit IBM Torino processor's hardware noise profile yielded critical insights for the Noisy Intermediate-Scale Quantum (NISQ) era. The empirical sweeps established that depolarizing and readout errors initially act as an analog to classical dropout, actively smoothing the decision boundary. However, the exact epoch thresholds captured during training prove that if execution continues indefinitely, the expressivity of the quantum circuit will detrimentally overfit to the physical hardware errors, cementing early stopping as a mandatory protocol.

To advance this research, future iterations of the benchmarking framework must continue to pivot toward quantum-native verification. The integration of the NTangled dataset in this thesis establishes a new standard for evaluating representational capacity against non-local, multipartite entangled states. Future work should focus on deploying the aligned QKA-ResQNet directly onto physical IBM quantum processors to validate the noise inflection thresholds identified in these simulations, solidifying the role of hybrid quantum networks in extreme biomedical and cheminformatics environments.